



Peeking Under The Hood

Low level visibility in Android with eBPF

Alex Calleja
Malware Researcher
[alex.calleja\(at\)zimperium.com](mailto:alex.calleja(at)zimperium.com)

BSides Malaga
14/3/2025



Outline

- Introduction to eBPF
- Motivation: eBPF for App analysis
- eBPF in Android
- Demo Time!
- eBPF Limitations
- Conclusions

whoami

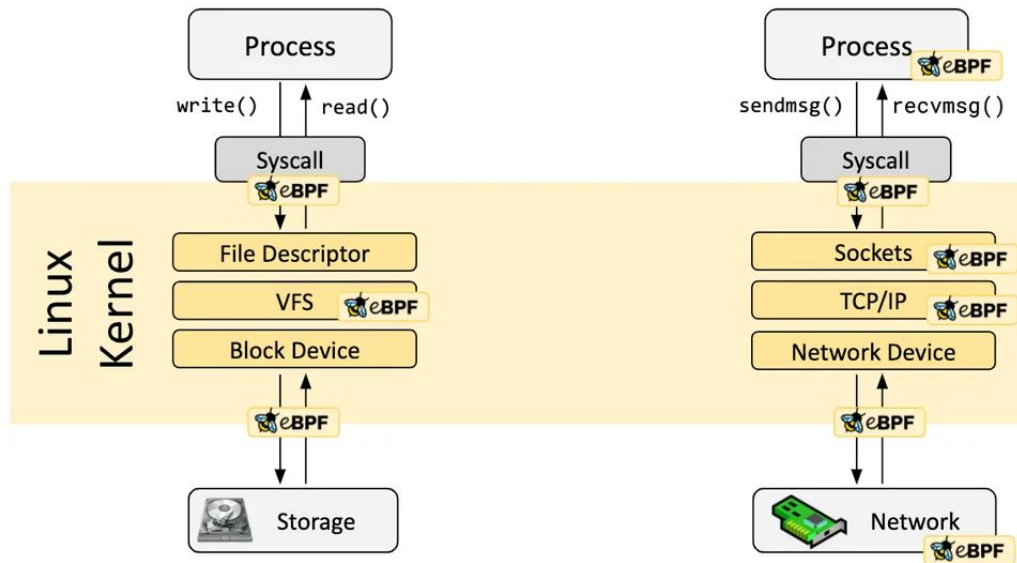
- Researcher @ Zimperium since 2018
- Android malware
- AOSP internals
- Leading the development of dynamic analysis platform
- Scaled the analysis pipeline to thousands of apps per day



`alex.calleja(at)zimperium.com`

Introduction to eBPF

- Instrument the kernel with event-driven programs
- Great observability due to privileged execution level
- Sandboxed environment & BPF validator ensure safety
- Applied across various fields, including networking, security, performance analysis, and more



Introduction to eBPF II

- BPF supports different program types with different functionalities
- Tracing programs allow attaching complex event handlers to system tracing probes

`/sys/fs/trace`

tracepoints

- Allows to instrument specific functions in the kernel, marked as traceable

k(ret)probes

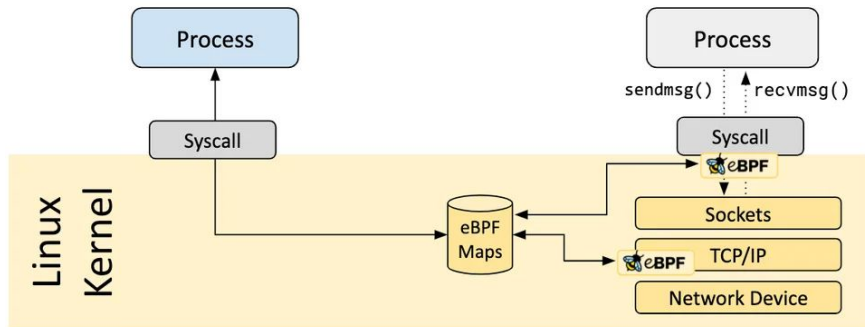

- Allows to intercept call/return events of kernel functions

u(ret)probes

- Allows to intercept call/return events of user-space code
- Requires specifying executable program and offset

Introduction to eBPF III

- Common setup consists in a **loader** + a **BPF implant**
- **CO-RE**: ensure compatibility across different kernel versions



- Different **data structures** allow communication between user-space and bpf side

Motivation: eBPF for App Analysis

Malware analysis

- Dynamic analysis often fails due to the detection of the emulated environment
- Frida usually saves the day, but falls short under specific circumstances
- Performance penalty when too much instrumentation is present

Motivation: eBPF for App Analysis

Packers/Runtime app self-protection reversing

- RASPs usually rely on inlined syscalls and other tricks (e.g hooked libc)
- Often include loops checking for root, emulator or dynamic binary instrumentation tooling
- This checks run usually on `.init` functions, which hinders further research

Motivation: Alternatives to eBPF

Dynamic binary Instrumentation (e.g Frida)

- 💪 Quick & easy solution
- 💪 Flexibility
- 💪 Decent compatibility across Android version

- 👎 Huge footprint
- 👎 Performance drops as agent size increases

Kernel Instrumentation (e.g Kernel modules)

- 💪 Full access to kernel APIs/types
- 💪 Full access to system & process resources
- 💪 Minimal overhead

- 👎 Requires painful kernel recompilation loops
- 👎 System crashes

BPF Programs

- ✌ Convenient trade-off between flexibility & observability

eBPF in Android

eBPF is well supported
in Android kernels

Since Android13 stock
kernels (5.10+) expose
BTF debugging info




Android






AOSP contains its own
version of libbpf
(libbpf_android.so)

AOSP includes BPF
implants for measuring
power & traffic

eBPF in Android

| Approach  | Pros  | Cons  |
|--|---|--|
| AOSP eBPF loader & examples | <ul style="list-style-type: none">• Native integration with Android APIs | <ul style="list-style-type: none">• Requires setting up AOSP build environment |
| NDK + Cilium | <ul style="list-style-type: none">• Good support for complex use-cases• High flexibility for advanced eBPF | <ul style="list-style-type: none">• Higher complexity• Requires minimum understanding of NDK and Go's build system |
| adeb + BCC | <ul style="list-style-type: none">• Ideal for experimental setup• Active community support for BCC | <ul style="list-style-type: none">• Performance drops due to BCC overhead• Dependency management |
| Build libbpf + loader from scratch | <ul style="list-style-type: none">• Full control over libbpf and program behavior• Highly flexible for specialized use-cases | <ul style="list-style-type: none">• High Integration effort• Limited community and support for Android-specific solutions |

eBPF in Android

| Approach  | Pros  | Cons  |
|--|--|--|
| AOSP eBPF loader & examples | <ul style="list-style-type: none">• Native integration with Android APIs | <ul style="list-style-type: none">• Requires setting up AOSP build environment |
| NDK + Cilium | <ul style="list-style-type: none">• Good support for complex use-cases• High flexibility for advanced eBPF | <ul style="list-style-type: none">• Higher complexity• Requires minimum understanding of NDK and Go's build system |
| adeb + BCC | <ul style="list-style-type: none">• Ideal for experimental setup• Active community support for BCC | <ul style="list-style-type: none">• Performance drops due to BCC overhead• Dependency management |
| Build <code>libbpf</code> + loader from scratch | <ul style="list-style-type: none">• Full control over <code>libbpf</code> and program behavior• Highly flexible for specialized use-cases | <ul style="list-style-type: none">• High Integration effort• Limited community and support for Android-specific solutions |

eBPF in Android: Building BPF programs from scratch

- An user space loader program must be written in order to load the bpf ELF file
- Alternatively, use `bpftool` and load `skel.h` file
- Linking loader against `libbpf` allows us to use `bpf_prog__*` API
- The loader must also contain the logic to consume the data sent from the implant or read from `trace_pipe`

```
struct bpf_object* obj;
struct bpf_program * prog;
obj = bpf_object__open_file("/data/local/tmp/implant.o.bpf", NULL);
/* load program, this issues a BPF(BPF_PROG_LOAD), > syscall */
if (bpf_object__load(obj)) {
    fprintf(stderr, "ERROR: loading BPF object file failed. Errno: %d\n");
    return 1;
}
/* find program in loaded ELF */
prog = bpf_object__find_program_by_name(obj, "handle_tp_sys_enter");
if (!prog) {
    fprintf(stderr, "ERROR: finding a prog in obj file failed\n");
    goto cleanup;
}
struct bpf_link *link = NULL;
/* associate program with tracepoint */
link = bpf_program__attach(prog);
```

```
void ringbuff_callback(void *ctx, int cpu, void *data, unsigned int data_sz){
    printf("New event received!\n")
}
int fd = bpf_object__find_map_fd_by_name(obj, "ringbuf_map");
if(fd<0){
    fprintf(stderr, "ERROR: could not find fd for map\n");
    return -1;
}
struct ring_buffer * rb = ring_buffer__new(fd, ringbuff_callback, NULL, NULL);
while (1) {
    err = ring_buffer__poll(rb, 100);
    if (err < 0) {
        printf("Error polling ring buffer: %d\n", err);
        break;
    }
}
```

eBPF in Android: Building BPF programs from scratch

- SEC macro defines sections within the resulting ELF file
- The program type defines the parameters that are accepted

implant: bpf_example.c

```
$(CLANG) -g -O2 -target bpf \
        -include vmlinux.h \
        -I../loader/dependencies/libbpf/src/ \
        I../loader/dependencies/libbpf/include/bpf \
        -c $^ -o $(IMPLANT_OBJ)
```

- Resulting bytecode can be checked with `llvm-objdump -S`


```
SEC("tp/raw_syscalls/sys_enter")
int handle_tp_sys_enter(struct trace_event_raw_sys_enter* ctx) {
    int pid = bpf_get_current_pid_tgid() >> 32;
    long syscall_nr = ctx->id;
    bpf_printk("PID %d issues syscall: %ld\n", pid, syscall_nr);
    return 0;
}
```

- Invoke clang with `-target bpf`

```
Disassembly of section tp/raw_syscalls/sys_enter:

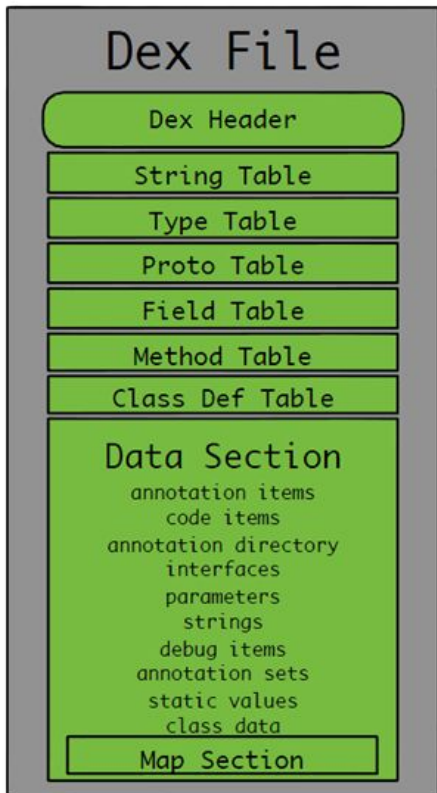
0000000000000000 <handle_tp_sys_enter>:
; int handle_tp_sys_enter(struct trace_event_raw_sys_enter* ctx) {
0:   bf 16 00 00 00 00 00 00 r6 = r1
;   int pid = bpf_get_current_pid_tgid() >> 32;
1:   85 00 00 00 0e 00 00 00 call 0xe
;   long syscall_nr = ctx->id;
2:   79 64 00 00 00 00 00 00 r4 = *(u64*)(r6 + 0x0)
;   int pid = bpf_get_current_pid_tgid() >> 32;
3:   77 00 00 00 20 00 00 00 r0 >>= 0x20
;   bpf_printk("PID %d issues syscall: %ld\n", pid, syscall_nr);
4:   18 01 00 00 00 00 00 00 00 00 00 00 00 00 00 00 r1 = 0x0
5:   b7 02 00 00 1d 00 00 00 r2 = 0x1d
6:   bf 03 00 00 00 00 00 00 r3 = r0
7:   85 00 00 00 06 00 00 00 call 0x6
8:   return 0;
9:   b7 00 00 00 00 00 00 00 r0 = 0x0
10:  95 00 00 00 00 00 00 00 exit
```

Demo Time!



- RootBeer can detect if the device has been rooted or tampered with
- The detection is mainly based on several file system checks (i.e `su`, `magisk` binaries)
- It also checks for specific property values by calling `getprop (ro.debuggable, ro.secure)` and grepping its output
- Additionally calls `mount` to check for `+rw` permissions on `system` partition
- **How can we evade this detections with `kprobes+eBPF`?**
 - Identify relevant syscalls (attach to `sys_enter tracepoint`)
 - Kprobe relevant syscalls
 - Use `bpf_probe_write_user` to smash arguments passed to the syscall

Demo Time!




- Often malware loads different DEX files during its execution
- This technique is used for hiding payloads and evade static analysis
- Frida can be used to instrument an unpacking point and dump the loaded DEX file
- What if the app is hardened against Frida?
- **Uprobes in action: dumping dex files without Frida**
 - Locate relevant function in user space (OpenCommon from libdexfile.so)
 - Place a Uprobe
 - Send back to user space {pid, dex_base, dex_size}
 - Use `process_vm_readv` to read `dex_size` bytes from target process memory

BPF Limitations

- Kernel memory is not writable
- Just a small subset of Kernel APIs/types are available
- The validator imposes several limitations such as:
 - Very restrictive pointer arithmetic
 - Loops must be deterministic (i.e it must be possible to check if a loop ends at load time)
- Program size restricted to 4096 instructions
- Despite efforts, compatibility is still an issue, mostly with older (<11) Android versions

Conclusions



- BPF provides a way to safely instrument the kernel with event-driven programs to gain observability of low level events
- It can be used for malware analysis and app analysis to overcome limitations of other approaches
- eBPF is well supported in Android kernels
- Still, some integration effort is required in Android
- Overall, eBPF is a valuable addition to your app analysis toolkit

Thank you!

Alex Calleja

Researcher

alex.calleja(at)zimperium.com

Code & slides available in github.com/alximw/BSidesMalaga25